

Exercise 3 — Build it with a durable working method

Stand-alone version of Exercise 3, for working on your own with **real GitHub Copilot custom agents** — not a simulation. The workspace bundled alongside this document already has the skill and the three agents configured as files; there is nothing to install or paste, only to open and confirm.

What you are building

The exact same domain as Exercise 2, from the exact same two specifications — nothing about the software changes. What changes is the *process*: instead of one long conversation, the way you use the assistant becomes something you can inspect, approve, and reuse, not just something that happened once in a chat window.

Before you start

1. Unzip the accompanying `exercise-3-copilot-workspace` folder and open **the whole folder** in VS Code — not a single file.
2. Confirm Java 21, Maven, Git, and your GitHub Copilot licence are available.
3. Open Copilot Chat and run **Chat: Open Customizations**. Confirm you can see:
 - `Specification Planner`, `Domain Implementer`, `Conformance Reviewer`;
 - `/create-plan`, `/implement-approved-plan`, `/verify-conformance`; and
 - the `implement-sell-stocks` skill.

If any item is missing, your Copilot licence or organisation policy does not currently support custom agents or skills — check that before continuing; nothing later in this document works without it.

The two specifications are already in `docs/spec/` inside the workspace — you do not need to place them yourself:

- `sell-stocks-spec.md` — behaviour: preconditions, the FIFO rule, the money definitions, and acceptance criteria AC-01 ... AC-24.
- `domain-class-diagram.puml` — structure: classes, fields, methods, visibility, and relationships.

Why this exercise exists

Exercise 2 made the **behaviour** durable: it lives in a spec file instead of your memory of a conversation. But the **process** you used to turn that spec into code — the questions you asked first, the order you

built things in, who checked the result and how — was still invisible. It happened once, inside one conversation, and vanished the moment the chat ended.

This exercise does for the *working method* what Exercise 2 did for the *domain*: it takes something that used to live only in a conversation and turns it into something you can read, review, and reuse — a **skill** for the repeatable method, and **role-scoped agents** for who is allowed to do what. Here, both are real, checked-in files: `.agents/skills/implement-sell-stocks/` and `.github/agents/*.agent.md`, already in the workspace you opened.

A concrete failure this catches

Picture Exercise 2 again. The specification never says what should happen at some edge case the acceptance criteria don't cover — say, a sale priced exactly at a lot's own purchase price. In one continuous conversation, the assistant can quietly pick an answer, write the code for it, then write a test that agrees with the code it just wrote — and you would very likely never notice, for the same reason Exercise 1's self-graded tests never failed: the actor that made the decision and the actor that checked it were the same conversation, with the same blind spot.

Exercise 3 puts two structural obstacles in that path, and here they are enforced by Copilot itself, not just by discipline:

1. **The Specification Planner agent has no edit tool.** It is configured to read and propose only — it cannot touch a file even if it tried. If it silently invents behaviour, that invention shows up as a line in `plan.md` you are explicitly asked to approve, not as a fait accompli buried in a diff you skim past.
2. **The Conformance Reviewer agent has no edit tool either**, and is selected as an independent pass over the finished implementation. It checks the code against the specification and the diagram directly, not against "does this look like what I would have written" — because Copilot never let it write anything in the first place.

Neither obstacle *guarantees* the assumption gets caught. What they guarantee is that it has to surface somewhere reviewable, instead of staying invisible inside one uninterrupted train of thought.

Exercise 2 vs. Exercise 3

Dimension	Exercise 2	Exercise 3
Domain inputs	Two specifications	The same two specifications
Process	One conversation, start to finish	Three role-scoped agents, invoked in order
Planning evidence	Whatever you remember from the chat	<code>plan.md</code> — written, dated, approved
Who checks the result	Often the same conversation that built it	The Conformance Reviewer agent, which has no edit tool at all
Reuse next time	Retype the prompt	The skill and agent files stay in the repository, discovered automatically

Is it worth the overhead? Be honest about this

At this kata's size, Exercise 3 is genuinely more work than Exercise 2's thirty-line prompt — three agent invocations and a written approval instead of one. That overhead does not pay for itself on a task this small, and "Worth trying afterwards" at the end asks you to feel that directly.

It starts paying off at a different scale than this kata: when a change is large or risky enough that you need to *prove*, after the fact, who approved what before it was built — closer to how a bank already treats a pull-request approval or a change record than it might first appear. It also pays off across time: a skill and a set of agents, written once, get reused on the next feature; a prompt typed once gets forgotten.

Hold both of those in mind as you work through the rest of this document — the four ideas below are what carries that value, not the paperwork itself.

The four ideas — and why each one matters

1. **A skill is a versioned, repeatable method** — not a one-off prompt you retype each time. It lives in `.agents/skills/implement-sell-stocks/SKILL.md`, reviewable like any other file in the repository, and Copilot loads it automatically when it's relevant.
2. **Role separation is a permission boundary, not a suggestion.** The Specification Planner has no edit tool at all, so it cannot quietly start implementing before you've approved anything; the Conformance Reviewer has no edit tool either, so it cannot "fix while reviewing" and blur checking into doing.
3. **A human checkpoint gates planning → implementation.** Nothing gets built until you've read the plan, changed at least one thing yourself, and written down that you approve it — turning approval into a real decision instead of a reflexive "looks good."
4. **Independent review happens before any fix.** The reviewer reports findings first; you decide what goes back to the implementer — keeping "did we build the right thing" separate from "do I like what got built."

Technical constraints

- Java 21, Maven, a standalone project with its own `pom.xml`, no parent.
- JUnit 5 only, in test scope. No Spring, no persistence, no REST layer.
- Package root: `com.neueda.portfolio.domain`.
- Every monetary amount as `BigDecimal`, scale 2, `HALF_UP` — never `double` or `float`.

These are already encoded in `AGENTS.md` and the agent files — you don't need to repeat them yourself; they're here so you know what to check for while reviewing.

Step 1 — Plan

Run this in Copilot Chat:

```
/create-plan
```

The Specification Planner reads both specifications but cannot edit your files — it has no edit tool. When it returns a proposed plan:

1. copy the proposal into `plan.md` ;
2. change or refine **at least one item yourself** — a split task, an added risk, a corrected dependency order, a criterion the proposal under-specified;
3. confirm AC-01 through AC-24 all appear in the coverage table;
4. change the plan status to `approved` ; and
5. fill in the Human checkpoint block: `Decision: approved` , your name, the date.

Do not continue until this file says `approved` in your own words, not the assistant's. This step is the point of the exercise — it is what makes approval a review act instead of a rubber stamp.

Step 2 — Implement

Run:

```
/implement-approved-plan
```

This generates `pom.xml` , `src/main` , and `src/test` following the approved plan. Review each checkpoint and proposed diff — do not approve everything automatically. Run `mvn test` yourself once it reports done; don't take "the tests pass" on trust from the transcript.

Step 3 — Independent review

Run:

```
/verify-conformance
```

The Conformance Reviewer has no edit tool, so it can only report findings, not fix them. Read the findings and classify each one: a verified defect, a suspected risk needing more evidence, an optional improvement outside the specification, or a false positive. Only findings **you** accept go back to the implementer.

Step 4 — Final check

Run this yourself from the workspace root:

```
mvn test
```

Expected result: exactly **36 tests**, no failures or errors. If Bash is available, you can also run `.agents/skills/implement-sell-stocks/scripts/verify-workspace.sh` for an additional mechanical check.

Completion contract

#	Requirement
1	<code>mvn test</code> passes with exactly 36 tests, no failures or errors
2	AC-01 through AC-24 are traceable in test names or display names
3	The implementation matches the class diagram — nothing extra, nothing missing
4	<code>plan.md</code> was approved, in your own words, before any implementation started
5	The Conformance Reviewer ran and reported before any fix
6	Only findings you accepted were sent back for correction
7	<code>plan.md</code> still accurately describes what was actually built

A green test run alone does not satisfy this contract — items 4 through 6 are about the *process*, and nothing in the code proves they happened.

Worth trying afterwards

- **Ask the Domain Implementer to also review its own work**, instead of switching to the Conformance Reviewer agent, and watch it agree with itself. It's the same blind spot as Exercise 1's self-written tests, wearing a different costume — an evaluator that already committed to an answer rarely reverses itself.
- **Approve a plan with a task deliberately left out**, then run `/implement-approved-plan`. Does it notice the gap and ask, or build past it silently?
- **Count the overhead**. Exercise 2 was one prompt of about thirty lines. This exercise was three agent invocations plus a written, human-approved checkpoint. Was that worth it at this kata's size? At what size would it be?

What's next

The domain became durable in Exercise 2. The working method became durable here. Exercise 4 keeps both and changes what they get pointed at: the same unchanged domain, now wrapped in a REST API and a real database.